

```
In [1]: import numpy as np
import pandas as pd
from scipy.stats import norm,t
import matplotlib.pyplot as plt
```

Q1

1.1

```
In [2]: def getClass1Prop(x, r):

    x = np.array(x)
    dist = np.zeros(len(x_train))
    for i in range(len(x_train)):
        dist[i] = np.linalg.norm(x-x_train[i])
    dist_label_1 = dist[y_train==1]
    dist_1r = dist_label_1[dist_label_1<=r]
    if len(dist[dist<=r]):
        return len(dist_1r)/len(dist[dist<=r])
    else:
        return np.nan
```

1.2

```
In [3]: def computeMisVal(data_val, r):
    x_val = data_val.iloc[:,1:3].to_numpy()
    y_val = data_val['Y'].to_numpy()
    y_pred = np.zeros(len(y_val))
    for i in range(len(x_val)):
        p = getClass1Prop(x_val[i],r)
        y_pred[i] = 1 if p>=0.5 else 0

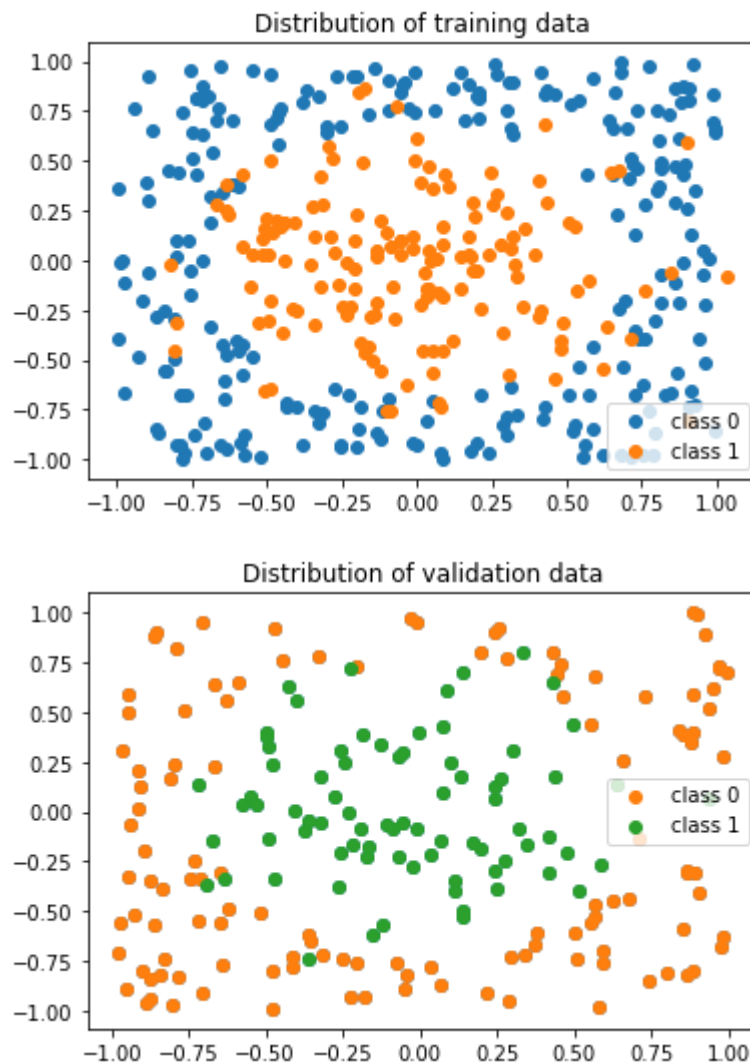
    mis = (len(y_val)-(y_pred==y_val).sum())/len(y_val)

    return mis
```

1.3

```
In [4]: data = pd.read_csv('SML.NN.data.csv')
data_train = data[data['set']=='train']
data_val = data[data['set']=='valid']
data_test= data[data['set']=='test']
x_train = data_train.iloc[:,1:3].to_numpy()
y_train = data_train['Y'].to_numpy()
x_val = data_val.iloc[:,1:3].to_numpy()
y_val = data_val['Y'].to_numpy()
```

```
In [5]: plt.figure()
plt.scatter(x_train[y_train==0][:,0],x_train[y_train==0][:,1],label = "class
0")
plt.scatter(x_train[y_train==1][:,0],x_train[y_train==1][:,1],label = "class
1")
plt.legend()
plt.title('Distribution of training data')
plt.show()
plt.figure()
plt.scatter(x_val[:,0],x_val[:,1])
plt.scatter(x_val[y_val==0][:,0],x_val[y_val==0][:,1],label = "class 0")
plt.scatter(x_val[y_val==1][:,0],x_val[y_val==1][:,1],label = "class 1")
plt.legend()
plt.title('Distribution of validation data')
plt.show()
```



The distributions of training and validation set show that points in class 1 mainly locate at center while points in class 0 form the outer circle. A good value of r should be able to differentiate two classes. The radius of points in class 1 is roughly 0.25. So $r = 0.25$ would be my guess.

1.4

```
In [6]: def computeMisTest(data_test, r):
        x_test = data_test.iloc[:,1:3].to_numpy()
        y_test = data_test['Y'].to_numpy()
        y_pred = np.zeros(len(y_test))
        for i in range(len(x_test)):
            p = getClass1Prop(x_test[i],r)
            y_pred[i] = 1 if p>=0.5 else 0


        mis = (len(y_test)-(y_pred==y_test).sum())/len(y_test)

        return mis
```

```
In [7]: data = pd.read_csv('SML.NN.data.csv')
        data_train = data[data['set']=='train']
        data_val = data[data['set']=='valid']
        data_test= data[data['set']=='test']
        x_train = data_train.iloc[:,1:3].to_numpy()
        y_train = data_train['Y'].to_numpy()
        r_grid = np.arange(0.01,1.01,0.01)
        mis_total = []
```

```
In [8]: for i in r_grid:
        mis_total.append(computeMisVal(data_val, i))

        plt.figure()
        plt.plot(r_grid,mis_total)
        plt.xlabel('r')
        plt.ylabel('misclassification rate')
        plt.show()
        rs = r_grid[np.argmin(mis_total)]
        print(f"When r = {rs}, the model achieves the lowest misclassification rate as {mis_total[np.argmin(mis_total)]}")
```



When $r = 0.26$, the model achieves the lowest misclassification rate as 0.06

```
In [9]: mis_test = computeMisTest(data_test, rs)
mis_test_gs = computeMisTest(data_test, 0.25)
print(f"When r = {rs}, the misclassification rate on test data is {mis_test}")
print(f"When r = {0.25}, the misclassification rate on test data is {mis_test_gs}")
```

When $r = 0.26$, the misclassification rate on test data is 0.045

When $r = 0.25$, the misclassification rate on test data is 0.05

0.25 is our guess value of r , 0.26 is the value chosen based on experiments. The results show our guess is very close to the optimal value of r

Q2

2.1

```
In [10]: from IPython.display import Image
Image('2_1.jpg')
```

Out[10]:

$$\begin{aligned}
 2.1 \quad f_Y(y) &= \int_{-\infty}^{\infty} f_{XY}(x, y) dx \\
 &= \frac{1}{2\pi b_X b_Y \sqrt{1-\rho^2}} \int_{-\infty}^{\infty} e^{-\frac{1}{2} Q(x, y)} dy \\
 \text{where } Q(x, y) &= \frac{\left(\frac{x-\mu_X}{b_X}\right)^2 - 2\rho\left(\frac{x-\mu_X}{b_X}\right)\left(\frac{y-\mu_Y}{b_Y}\right) + \left(\frac{y-\mu_Y}{b_Y}\right)^2}{1-\rho^2} \\
 &= \left(\frac{x-a}{b}\right)^2 + c \\
 a &= \mu_X + \rho \frac{b_X}{b_Y} (y - \mu_Y) \\
 b &= b_X \sqrt{1-\rho^2} \\
 c &= \left(\frac{y-\mu_Y}{b_Y}\right)^2 \\
 \Rightarrow f_Y(y) &= \frac{1}{2\pi b_X b_Y \sqrt{1-\rho^2}} \int_{-\infty}^{\infty} e^{-\frac{1}{2} \left[\left(\frac{x-a}{b}\right)^2 + c\right]} dx \\
 &= \frac{e^{-\frac{c}{2}}}{2\pi b_X b_Y \sqrt{1-\rho^2}} \int_{-\infty}^{\infty} e^{-\frac{1}{2} \left(\frac{x-a}{b}\right)^2} dx \\
 &= \frac{e^{-\frac{c}{2}}}{\sqrt{2\pi} b_X b_Y \sqrt{1-\rho^2}} b_X \sqrt{1-\rho^2} \\
 &= \frac{1}{\sqrt{2\pi} b_Y} e^{-\frac{1}{2} \left(\frac{y-\mu_Y}{b_Y}\right)^2}
 \end{aligned}$$



Similarly,

$$\begin{aligned}
 f_X(x) &= \int_{-\infty}^{\infty} f_{XY}(x, y) dy \\
 &= \frac{1}{\sqrt{2\pi} b_X} e^{-\frac{1}{2} \left(\frac{x-\mu_X}{b_X}\right)^2}
 \end{aligned}$$

```
In [11]: x = np.arange(-3, 3, 0.01)
y = np.arange(-3, 3, 0.01)

# Calculating mean and standard deviation
mu_x,mu_y,sig_x,sig_y = 0,0,1,1

plt.figure()
plt.plot(x, norm.pdf(x, mu_x, sig_x))
plt.title('Marginal PDF of x')
plt.show()
plt.figure()
plt.plot(y, norm.pdf(y, mu_y, sig_y))
plt.title('Marginal PDF of y')
plt.show()
```



Those marginal pdfs are identical with Normal(0,1) pdf

2.2

In [12]: Image('2_2.jpg')

Out[12]:

2.2 If $\rho = 0$,

$$f_{XY}(x, y) = \frac{1}{2\pi b_x b_y} e^{-\frac{1}{2} \left[\left(\frac{x - \mu_x}{b_x} \right)^2 + \left(\frac{y - \mu_y}{b_y} \right)^2 \right]}$$

$$f_X(x) f_Y(y) = \frac{1}{2\pi b_x b_y} e^{-\frac{1}{2} \left[\left(\frac{x - \mu_x}{b_x} \right)^2 + \left(\frac{y - \mu_y}{b_y} \right)^2 \right]}$$

$$\therefore f_{XY}(x, y) = f_X(x) f_Y(y)$$

$\therefore X$ and Y are independent


2.3

Given $Y = 1$:

$$E[X|Y=1] = \mu_x + \rho \sigma_x \frac{1 - \mu_y}{\sigma_y}$$

$$\text{Var}[X|Y = 1] = (1 - \rho^2) \sigma_x^2$$

```
In [13]: rho_grid = np.arange(-0.75,1,0.25)
x = np.arange(-5, 5, 0.01)
mu_x,mu_y,sig_x,sig_y = 0,0,1,1
plt.figure()
for rho in rho_grid:
    mu = mu_x+rho*sig_x*(1-mu_y)/sig_y
    sig = (1-rho**2)*sig_x**2
    plt.plot(x, norm.pdf(x, mu, sig),label=f'rho = {rho}')
plt.legend()
plt.show()
```



The negative and positive ρ both help to predict X since they indicate a linear relation between X and Y . Larger ρ implies a stronger linearity. It is hard to predict X if X and Y are independent, i.e. $\rho = 0$.

Q3

3.1

```
In [14]: np.random.seed(0)
lam = 10
x = np.random.exponential(1/lam,100)
```

$$L = \lambda e^{(-\lambda x_1)} \lambda e^{(-\lambda x_1)} \dots \lambda e^{(-\lambda x_{100})} = \lambda^{100} e^{(-\lambda(x_1 + x_2 + \dots + x_{100}))}$$

$$L' = \ln L = 100 \ln \lambda - \lambda \sum_{i=0}^{100} x_i$$

$$\frac{\partial L'}{\partial \lambda} = \frac{100}{\lambda} - \sum_{i=0}^{100} x_i = 0$$

$$\lambda = \frac{100}{\sum_{i=0}^{100} x_i}$$

```
In [15]: lam = 100/x.sum()
print(f'estimated lambda is {lam}')
```

estimated lambda is 10.885557365453474

The estimated value of λ (10.89) is close to true λ (10)

3.2

In [16]: Image('3_2.jpg')

Out[16]:

$$3.2 \quad f(x) = \begin{cases} \lambda e^{-\lambda x} & x \geq 0 \\ 0 & x < 0 \end{cases}$$

$$L(\lambda) = \prod_{i=1}^N \lambda e^{-\lambda x_i}$$

Take \ln :

$$\begin{aligned} \Rightarrow L' = \ln L &= \sum_{i=1}^N \ln \lambda e^{-\lambda x_i} \\ &= \sum_{i=1}^N (\ln \lambda - \lambda x_i) \\ &= N \ln \lambda - \lambda \sum_{i=1}^N x_i \end{aligned}$$

Take derivative:

$$\Rightarrow \frac{\partial L'}{\partial \lambda} = \frac{N}{\lambda} - \sum_{i=1}^N x_i = 0$$

$$\Rightarrow \lambda = \frac{N}{\sum_{i=1}^N x_i}$$

Check if the solution is the maximum:

$$\frac{\partial^2 L'}{\partial \lambda^2} = -\frac{N}{\lambda^2}$$

$$N > 0, \lambda^2 > 0$$

$$\therefore \frac{\partial^2 L'}{\partial \lambda^2} = -\frac{N}{\lambda^2} < 0$$

$$\therefore \lambda = \frac{N}{\sum_{i=1}^N x_i}, \text{ which is the reciprocal of sample mean}$$

Q4

Case 1

```
In [17]: seed_grid = np.arange(1,1001)
coverage = 0

for i in seed_grid:
    np.random.seed(i)
    x = np.random.normal(0,1,4)
    mu = x.mean()
    std = x.std()
    up = mu+1.6449*std/np.sqrt(4)
    lower = mu-1.6449*std/np.sqrt(4)
    if up>=0 and lower<=0:
        coverage += 1
    else:
        pass
ef = coverage/len(seed_grid)
print(f'The empirical frequency of coverage is {ef}')
```

The empirical frequency of coverage is 0.749


Case 2

```
In [18]: seed_grid = np.arange(1,1001)
coverage = 0
width = []

for i in seed_grid:
    np.random.seed(i)
    x = np.random.normal(0,1,4)
    mu = x.mean()
    std = x.std()
    interval = t.interval(alpha=0.95, df=3, loc=mu, scale=std)
    if interval[1]>=0 and interval[0]<=0:
        coverage += 1
    else:
        pass
    width.append(interval[1]-interval[0])
ef = coverage/len(seed_grid)
print(f'The empirical frequency of coverage is {ef}')
```

The empirical frequency of coverage is 0.988

```
In [19]: plt.figure()
plt.hist(width)
plt.title('Histogram of Interval Widths')
plt.show()
```




Case 3

```
In [20]: seed_grid = np.arange(1,1001)
coverage = 0
width = []

for i in seed_grid:
    np.random.seed(i)
    x = np.random.normal(0,1,4)
    mu = x.mean()
    std = x.std()
    interval = norm.interval(alpha=0.95, loc=mu, scale=std)
    if interval[1]>=0 and interval[0]<=0:
        coverage += 1
    else:
        pass
    width.append(interval[1]-interval[0])
ef = coverage/len(seed_grid)
print(f'The empirical frequency of coverage is {ef}')
```

The empirical frequency of coverage is 0.957

```
In [21]: plt.figure()  
plt.hist(width)  
plt.title('Histogram of Interval Widths')  
plt.show()
```



Case 2 has the highest empirical frequency of coverage when exact small sample CI is computed using t-distribution. That makes sense since the sample is very small ($N=4$). When using approximate large-sample CI in case 3, the empirical frequency decreases.

```
In [ ]:
```