

STA6703 SML HW4

Christopher and Byron

Load libraries

```
library(MASS)
library(ISLR2)
```

```
##
## Attaching package: 'ISLR2'

## The following object is masked from 'package:MASS':
##
## Boston
```

Define a utility function

```
MCR <- function(true_vals, pred_probs, threshold=0.5){
  if(length(true_vals)!=length(pred_probs)){
    print("ERROR: predictions and true values not of same shape")
  }else{
    pred_vals = as.integer((pred_probs > threshold))
    mcr = sum(pred_vals != true_vals)/length(true_vals)
    return(mcr)
  }
}
```

Chapter 4

Question 5

5.a

LDA is better on the test set and QDA is better with the training set. QDA is able to describe non-linear boundaries and LDA only able to describe linear boundaries. So if the test set has a linear boundary the LDA technique will generalize better, but the QDA technique will over fit to the training data easier.

5.b

QDA will be better in the training set and in the testing set.

5.c

With more data QDA will become better at estimating the true boundary. With more data we expect the sample to be a more accurate representation of the test data. Therefore the QDA method will generalize better to the test data. The effect of over fitting will decrease.

5.d

False, LDA will better fit a linear decision boundary. QDA could provide an over-fitting model that will perform better on the training set, but worse on the test set. LDA will probably fit the linear decision boundary better than QDA and result in a lower test error rate.

Question 9

9.a

$$\begin{aligned} Odds &= \frac{P(X)}{1 - P(X)} \\ 0.37 &= \frac{P(X)}{1 - P(X)} \\ P(X) &= \frac{0.37}{1.37} = 0.27 \end{aligned}$$

With an odds of 0.37 it means 27% of people will default on their credit card payments.

9.b

$$\begin{aligned} Odds &= \frac{0.16}{1 - 0.16} \\ Odds &= 0.19 \end{aligned}$$

With a 16% chance of defaulting, the odds that she will default is 0.19.

Question 11

11.a

```
auto_data = Auto
mpg_med = median(auto_data$mpg); mpg_med
```


```
## [1] 22.75
```

```
auto_data$mpg01 = as.integer((auto_data$mpg > mpg_med)); set.seed(3)
auto_data[sample(1:length(auto_data[, 1]), 5, replace = FALSE), c(1, 9, 10)]
```

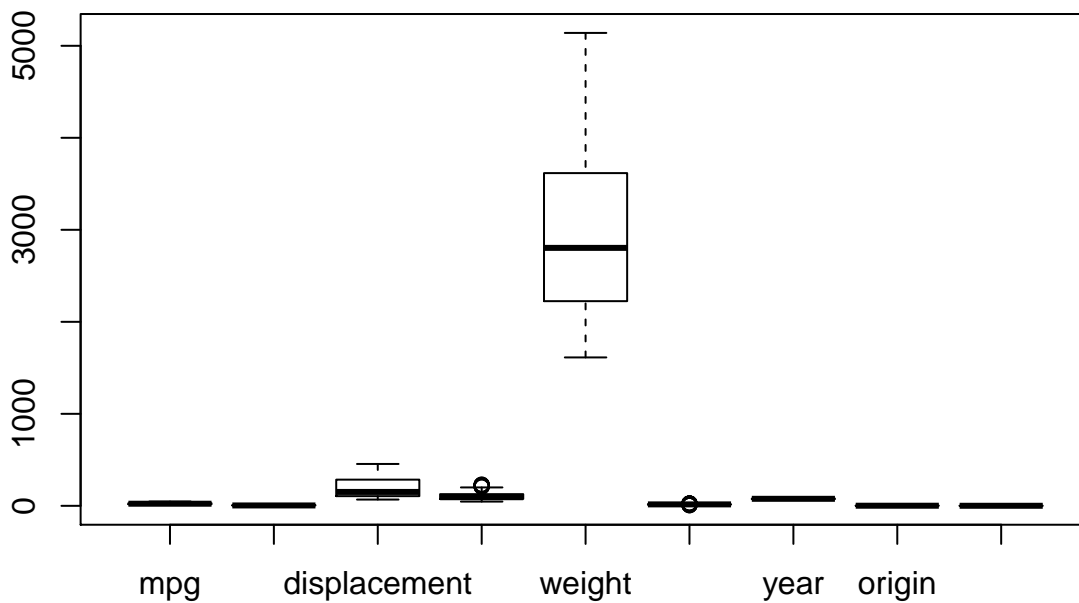
```
##      mpg                                name mpg01
## 263 19.2      chevrolet monte carlo landau      0
## 188 17.5 chevrolet chevelle malibu classic      0
## 142 29.0                                audi fox      1
## 37  19.0                                ford torino 500      0
## 396 28.0                                ford ranger      1
```

11.b

```
# heatmap(cor(auto_data[, -9]), Colv = NA, Rowv = NA, revC = TRUE)
library(pheatmap)
pheatmap(cor(auto_data[, -9]), display_numbers = T)
```



```
boxplot(auto_data[, c(-9)])
```



Use cylinders, displacement, weight and, horsepower as they all have strong negative correlations with mpg01. mpg has a strong positive correlation, but was used to create mpg01 so it is not independent.

11.c

```
set.seed(42)
sample <- sample.int(n = nrow(auto_data),
                    size = floor(0.75*nrow(auto_data)),
                    replace = F)

train <- auto_data[sample, ]
test  <- auto_data[-sample, ]
```

11.d

```
lda_auto = lda(mpg01 ~ cylinders+displacement+horsepower+weight, data=train)
lda_auto
```

Call:

```
## lda(mpg01 ~ cylinders + displacement + horsepower + weight, data = train)
##
## Prior probabilities of groups:
##      0      1
## 0.5340136 0.4659864
##
## Group means:
##  cylinders displacement horsepower  weight
## 0  6.777070      272.8599  129.44586 3598.854
## 1  4.218978      117.4562   78.51095 2345.971
##
## Coefficients of linear discriminants:
##                      LD1
## cylinders      -0.3763366347
## displacement -0.0018979140
## horsepower      0.0004421014
## weight        -0.0008794090
```

```
lda_auto_probs = data.frame(
  predict(lda_auto,
    test)$posterior[,2]
)
```

```
MCR(
  true_vals=test$mpg01,
  pred_probs=lda_auto_probs[,1],
  threshold=0.5)
```

```
## [1] 0.07142857
```

11.e

```
qda_auto = qda(mpg01 ~ cylinders+displacement+horsepower+weight, data=train)
qda_auto
```

```
## Call:
## qda(mpg01 ~ cylinders + displacement + horsepower + weight, data = train)
##
## Prior probabilities of groups:
##      0      1
## 0.5340136 0.4659864
##
## Group means:
##  cylinders displacement horsepower  weight
## 0  6.777070      272.8599  129.44586 3598.854
## 1  4.218978      117.4562   78.51095 2345.971
```

```
qda_auto_probs = data.frame(
  predict(qda_auto,
    test)$posterior[,2]
)
```

```
MCR(
  true_vals=test$mpg01,
  pred_probs=qda_auto_probs[,1],
  threshold=0.5)
```

```
## [1] 0.08163265
```

11.f

```
lr_auto = glm(mpg01 ~ cylinders+displacement+horsepower+weight,
  data=train,
  family="binomial")
```

```
summary(lr_auto)
```

```
##
## Call:
## glm(formula = mpg01 ~ cylinders + displacement + horsepower +
##      weight, family = "binomial", data = train)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -2.2583  -0.2785  -0.0115   0.3798   3.2536
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)  10.9255004  1.8727915   5.834 5.42e-09 ***
## cylinders    -0.0020851  0.3750310  -0.006  0.99556
## displacement -0.0115341  0.0089275  -1.292  0.19636
## horsepower   -0.0465546  0.0147880  -3.148  0.00164 **
## weight       -0.0016258  0.0007472  -2.176  0.02956 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 406.21  on 293  degrees of freedom
## Residual deviance: 165.91  on 289  degrees of freedom
## AIC: 175.91
##
## Number of Fisher Scoring iterations: 7
```

```
lr_auto_probs = data.frame(  
  predict(lr_auto,  
    test,  
    type = "response"  
  )  
)
```

```
MCR(  
  true_vals=test$mpg01,  
  pred_probs=lr_auto_probs[,1],  
  threshold=0.5)
```

```
## [1] 0.07142857
```

Typed Problem 1

```
# Read and give additional format to the data
data <- read.fwf("challenger/challenger.dat", width=c(8,8,8,8),
               col.names=c("flight", "degrees", "numfail", "indfail"))[1:23, ]

y = matrix(data$indfail, ncol = 1)
x = matrix(data$degrees, ncol = 1)
X = matrix(c(x^0, x), ncol = 2)
```

1.a

```
logLikelihood = function(b) {
  s = X %*% b
  p = exp(s) / (1 + exp(s))
  # logpy = sum(log(p) * (y == 1)) + sum(log(1 - p) * (y == 0))
  logpy = sum(dbinom(y, 1, p, log = TRUE))
  return(-logpy)
}
```

1.b

```
init = matrix(c(0, 0), ncol = 1)
mle = optim(init, logLikelihood, method = "BFGS")$par
c("(Intercept)" = mle[1], "degrees" = mle[2])
```

```
## (Intercept)      degrees
## 15.0940833 -0.2329155
```

```
glm1 = glm(indfail ~ degrees, family = binomial, data = data)
glm1$coefficients
```

```
## (Intercept)      degrees
## 15.0429016 -0.2321627
```



```
summary(glm1)
```

```
##
## Call:
## glm(formula = indfail ~ degrees, family = binomial, data = data)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -1.0611  -0.7613  -0.3783   0.4524   2.2175
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)  15.0429     7.3786   2.039  0.0415 *
## degrees      -0.2322     0.1082  -2.145  0.0320 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 28.267  on 22  degrees of freedom
## Residual deviance: 20.315  on 21  degrees of freedom
## AIC: 24.315
##
## Number of Fisher Scoring iterations: 5
```

As we can see, the obtained values in both cases are similar, the deviations could be due to the difference on the numerical methods used in both cases.

1.c

(i) First, a function to get the Hessian of function f

```
delta = 1e-4
gradF = function(f, x) {
  p = length(x)
  X = t(matrix(rep(x, p), ncol = p))
  J = (f(X + delta * diag(p)) - f(X - delta * diag(p))) / (2 * delta)
  basis = diag(p)
  for (i in 1:p){
    J[i] = (f(x + delta * basis[, i]) - f(x - delta * basis[, i])) / (2 * delta)
  }
  return(t(J))
}
```

```

hessF = function(f, x) {
  g = function(x) gradF(f, x)
  p = length(x)
  H = diag(p)
  basis = diag(p)
  for (i in 1:p){
    H[, i] = (g(x + delta * basis[, i]) - g(x - delta * basis[, i])) / (2 * delta)
  }
  return(H)
}
# test:
f = function(x) x[1]^2 + 2 * x[2]^2 - 2 * x[1] * x[2] + exp(x[1])
gradF(f, matrix(c(0, 0), ncol = 1))

```

```

##      [,1] [,2]
## [1,]    1    0

```

```

hessF(f, matrix(c(0, 0), ncol = 1))

```

```

##      [,1] [,2]
## [1,]    3   -2
## [2,]   -2    4

```

(ii) Approximate Hessian of the log-Likelihood

```

H = hessF(logLikelihood, mle); H

```

```

##      [,1]      [,2]
## [1,]  3.257413 221.4379
## [2,] 221.437871 15138.1628

```

(iii) Approximate the covariance matrix

```

S = solve(H); S

```

```

##      [,1]      [,2]
## [1,] 54.7458259 -0.80081046
## [2,] -0.8008105  0.01178015

```

```

summary(glm1)$cov.unscaled

```

```

##      (Intercept)      degrees
## (Intercept)  54.4441826 -0.79638547
## degrees     -0.7963855  0.01171512

```

(iv) Approximate the standar errors

```
sde = sqrt(diag(S)); c("(Intercept)" = sde[1], "degrees" = sde[2])
```

```
## (Intercept)      degrees  
##    7.3990422    0.1085364
```

```
summary(glm1)$coefficients[, 2] # second column (Std. Error)
```

```
## (Intercept)      degrees  
##    7.3786301    0.1082364
```

As expected, the approximation of the standard deviation using the estimated covariance matrix, is similar to the standard errors found by `glm`.


Problem 2

Import data

```
# load data
data <- read.csv("SML.NN.data.csv")
train_data = data[data$set == 'train' | data$set == 'valid',]
test_data = data[data$set == 'test',]
```

```
plot(train_data$X1,
     train_data$X2,
     pch=8,
     col=factor(train_data$Y),
     main='Training Data',
     xlab="X1",
     ylab="X2")
```

```
legend(1.3,
      1.5,
      legend=c('1', '0'),
      col=c('r', 'b'),
      fill=2:1,
      bg="white")
```




```

plot(test_data$X1,
     test_data$X2,
     pch=8,
     col=factor(test_data$Y),
     main='Testing Data',
     xlab="X1",
     ylab="X2")

legend(1.3,
      1.5,
      legend=c('1', '0'),
      col=c('r', 'b'),
      fill=2:1,
      bg="white")

```



Train models

L1

```

L1 = glm(Y ~ 1 + X1 + X2,
         data=train_data,
         family="binomial")

```

```
summary(L1)
```

```
##
## Call:
## glm(formula = Y ~ 1 + X1 + X2, family = "binomial", data = train_data)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -1.0497  -0.9987  -0.9498   1.3715   1.4516
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept) -0.46905     0.08403  -5.582 2.38e-08 ***
## X1          -0.12055     0.14660  -0.822   0.411
## X2           0.05727     0.14406   0.398   0.691
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 799.75  on 599  degrees of freedom
## Residual deviance: 798.96  on 597  degrees of freedom
## AIC: 804.96
##
## Number of Fisher Scoring iterations: 4
```

L2

```
L2 = glm(Y ~ 1 + X1 + X2 + X1^2 + X2^2 + X1*X2,
         data=train_data,
         family="binomial")
```

```
summary(L2)
```

```
##
## Call:
## glm(formula = Y ~ 1 + X1 + X2 + X1^2 + X2^2 + X1 * X2, family = "binomial",
##      data = train_data)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -1.2289  -0.9986  -0.8985   1.3708   1.5306
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept) -0.45827     0.08440  -5.430 5.65e-08 ***
## X1          -0.12868     0.14745  -0.873   0.383
```

```
## X2          0.05857    0.14474    0.405    0.686
## X1:X2       -0.49418    0.25100   -1.969    0.049 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
## Null deviance: 799.75  on 599  degrees of freedom
## Residual deviance: 795.04  on 596  degrees of freedom
## AIC: 803.04
##
## Number of Fisher Scoring iterations: 4
```

LDA

```
D1 = lda(Y ~ X1 + X2,
         data=train_data)
```

```
D1
```

```
## Call:
## lda(Y ~ X1 + X2, data = train_data)
##
## Prior probabilities of groups:
##      0      1
## 0.615 0.385
##
## Group means:
##           X1           X2
## 0  0.004605339 -0.02707597
## 1 -0.033643963 -0.01087806
##
## Coefficients of linear discriminants:
##           LD1
## X1 -1.6168003
## X2  0.7681857
```

QDA

```
D2 = qda(Y ~ X1 + X2,
         data=train_data)
```

```
D2
```

```
## Call:
## qda(Y ~ X1 + X2, data = train_data)
```

```
##
## Prior probabilities of groups:
##      0      1
## 0.615 0.385
##
## Group means:
##           X1           X2
## 0  0.004605339 -0.02707597
## 1 -0.033643963 -0.01087806
```

Test models

L1

```
L1_probs = data.frame(
  predict(L1,
    test_data,
    type = "response"
  )
)

MCR(
  true_vals=test_data$Y,
  pred_probs=L1_probs[,1],
  threshold=0.5)
```

```
## [1] 0.325
```

L2

```
L2_probs = data.frame(
  predict(L2,
    test_data,
    type = "response"
  )
)

MCR(
  true_vals=test_data$Y,
  pred_probs=L2_probs[,1],
  threshold=0.5)
```

```
## [1] 0.335
```


LDA

```
D1_probs = data.frame(  
  predict(D1,  
    test_data)$posterior[,2]  
)  
  
MCR(  
  true_vals=test_data$Y,  
  pred_probs=D1_probs[,1],  
  threshold=0.5)
```

```
## [1] 0.325
```

QDA

```
D2_probs = data.frame(  
  predict(D2,  
    test_data)$posterior[,2]  
)  
  
MCR(  
  true_vals=test_data$Y,  
  pred_probs=D2_probs[,1],  
  threshold=0.5)
```

```
## [1] 0.09
```

Visualize decision boundaries

```
# define threshold  
threshold = 0.5  
  
# create grid of points  
axis_ticks = seq(-2,2,0.1)  
grid_df = expand.grid(axis_ticks,axis_ticks)  
colnames(grid_df) <- c("X1","X2")  
  
# L1 prediction  
L1_probs = data.frame(  
  predict(L1,  
    grid_df,  
    type ="response"  
)
```

```

    )

grid_df$L1 = as.integer((L1_probs>threshold))

# L2 prediction
L2_probs = data.frame(
    predict(L2,
            grid_df,
            type ="response"
    )
)

grid_df$L2 = as.integer((L2_probs>threshold))

# D1 prediction
D1_probs = data.frame(
    predict(D1,
            grid_df)$posterior[,2]
    )

grid_df$D1 = as.integer((D1_probs>threshold))

# D2 prediction
D2_probs = data.frame(
    predict(D2,
            grid_df)$posterior[,2]
    )

grid_df$D2 = as.integer((D2_probs>threshold))

```

L1


```

plot(grid_df$X1,
     grid_df$X2,
     pch=8,
     col=factor(grid_df$L1),
     main='L1',
     xlab="X1",
     ylab="X2")

legend(1.3,
      1.5,
      legend=c('1', '0'),
      col=c('r', 'b'),
      fill=2:1,
      bg="white")

```

L1




L2

```
plot(grid_df$X1,
     grid_df$X2,
     pch=8,
     col=factor(grid_df$L2),
     main='L2',
     xlab="X1",
     ylab="X2")

legend(1.3,
      1.5,
      legend=c('1', '0'),
      col=c('r', 'b'),
      fill=2:1,
      bg="white")
```

L2




LDA

```
plot(grid_df$X1,
     grid_df$X2,
     pch=8,
     col=factor(grid_df$D1),
     main='D1',
     xlab="X1",
     ylab="X2")

legend(1.3,
      1.5,
      legend=c('1', '0'),
      col=c('r', 'b'),
      fill=2:1,
      bg="white")
```

D1



QDA

```
plot(grid_df$X1,
     grid_df$X2,
     pch=8,
     col=factor(grid_df$D2),
     main='D2',
     xlab="X1",
     ylab="X2")

legend(1.3,
      1.5,
      legend=c('1', '0'),
      col=c('r', 'b'),
      fill=2:1,
      bg="white")
```

D2

