

# ABE6933 SML HW3

## Directions

Please submit **one PDF** file including all your reports (answer + code + figures + comments; must be easily readable and have file size under a few megabytes) and **one R code script**. The R script is supplementary material to ensure that your code runs correctly. If you are using RMarkdown, please also include your `.Rmd` file.

Place these two (or three) files in a folder, make a zip or rar archive, and submit the archive electronically via Dropbox file request at [tinyurl.com/nbliznyuk-submit-files](https://tinyurl.com/nbliznyuk-submit-files) (on the landing page, enter your name so that we know it is you and email so that you get a confirmation).

For the full list of rules/policies/expectations, please visit “hw.rules.pdf” document.

**Deadline:** 08-Oct-2020, 11:59 PM EST.

## Practice/Optional Problems (do not submit)

1. Use your solution for the required Typed Problem 1 below to devise an iterative procedure to compute the rank of an  $n \times p$  matrix  $A$ . Verify your answer using the SVD decomposition (count the number of singular values numerically different from 0; `svd` function in R) or by QR decomposition (`qr` function in R).
2. Explore the use of multiple linear regression/`lm()` function to "solve" general linear systems  $Ax = b$ , where  $A$  is a general  $n \times p$  matrix.
3. Suppose one is given a matrix  $A$  of size  $n \times p$ ,  $n > p$ . Use `lm()` function to write your own routine to perform carry out the Gram-Schmidt orthogonalization process thereby obtaining a crude way to the QR decomposition of  $A$ .
4. Solving square linear systems  $Ax = b$  using LU and QR factorizations when  $A$  is invertible:

- (a) The LU factorization/decomposition is available as a part of the `Matrix` library in R. Mathematically,  $A = P \cdot L \cdot U$ , where  $P$  is a permutation matrix,  $L$  is a lower-triangular matrix and  $U$  is an upper-triangular matrix; all matrices are square.

Given the  $LU$  factorization of  $A$ , show how to solve the original linear system  $Ax = b$ .

- (b) The QR factorization/decomposition is available as a part of the base library in R. For a square  $A$  of full rank,  $A = Q \cdot R$ , where  $Q$  is a matrix with orthonormal columns (i.e.,  $Q'Q = QQ' = I$ ) and  $R$  is upper-triangular; both matrices are square.

Given the  $QR$  factorization of  $A$ , show how to solve the original linear system  $Ax = b$ .

*Hint to both subproblems:* Determine how to solve lower- and upper-triangular square systems by hand, i.e.,  $Ly = c$  and  $Uz = d$ . Start with a  $2 \times 2$  case, then generalize.

## Required Problems (for submission)

**Typed Problem 1.** Let  $X$  be an  $n \times p$  design matrix with  $n > p$  (predictors are in the columns); assume that the column of ones is the first column of  $X$ . Argue/show that, if columns of  $X$  are not linearly independent, then regressing the  $i$ th column of  $X$  on all other columns (for  $i = 2, \dots, p$ ) will identify this.

*Hint:* if  $R^2 = 1$ , what is the residual sum of squares (RSS)? Now relate this to the definition of linear dependence of columns of a matrix. It may help you build intuition if you generate such a matrix and carry out the proposed procedure before providing a conceptual answer.

**Typed Problem 2.** Verify or refute the following statement: if numerical predictors/covariates (i.e., columns of the design matrix  $X$ ) are linearly independent, then they are uncorrelated.

*Hint 1:* make sure you entirely understand the two definitions.

*Hint 2:* this problem does NOT assume you know how to “prove” mathematical statements.

**Typed Problem 3.** Multiple linear regression (and its flavors) implemented “by hand” in R.

3.1. Implement by hand an R function `myOLS` following the interface below.

You are not allowed to call other functions that do statistics (e.g., no `lm()` calls); you should use linear algebra operations (such as `solve(A)` to find the inverse of  $A$ ).

```
myOLS <- function(Y, X, is1 = TRUE) {  
  # Inputs:  
  # * Y is the vector of length n of response variables  
  # * X is an n-by-p matrix of numerical covariates (in columns); p < n  
  # ** assume the columns of X are linearly independent and  
  # ** do not include the column for the intercept as a part of the X matrix  
  # * is1 is a logical "flag" whether the intercept is included; is1 = TRUE by default  
  # Output:  
  # the function must return a list L with two elements:  
  # L[1] will contain the vector of OLS/MLE coefficients, betahat  
  # L[2] will contain standard errors (i.e., estimated standard deviations) for betahat  
}
```

Compare the results with those produced by `lm()` on the following simulated data:

```
n = 30; set.seed(0); p = 3;  
X = matrix(runif(n*p), nrow=n)*2-1;  
b = seq(1,p,by=1);  
Y = X%*%b + rnorm(n);  
fit1 = lm(Y ~ X); summary(fit1); # regression with an intercept  
fit0 = lm(Y ~ -1 + X); summary(fit0) # regression without an intercept
```

3.2. Use your implementation of `myOLS` to implement polynomial regression with one covariate, i.e.,  $Y = b_0 + b_1x + \dots + b_kx^k + \epsilon$ . Intercept is always included. The interface is below.

```
myPolyReg1 <- function(Y, X1, deg=1) {  
  # Inputs: same as for myOLS, except  
  # * X1 is a vector of length n that contain the covariate values (numerical)  
  # * deg is the degree k (i.e., largest power) of the polynomial fit; k < n ; deg=1 by default.  
  # Outputs: same as for myOLS  
}
```

Compare the results with those produced by `lm()` on the following simulated data:

```
n = 30; set.seed(0);
X = runif(n)*4-2; # X is uniformly distributed on [-2,2]
Y = 1 + 3*X - 2*X^2 + 1*X^3 + rnorm(n);
fit0 = lm(Y ~ X + I(X^2) + I(X^3)); summary(fit0)
```

3.3. Use your implementation of `myOLS` to implement a one-way ANOVA model, i.e., regression with a single categorical covariate. The interface for the function is below.

```
myAnova1 <- function(Y, XF, is1=TRUE) {
# Inputs: same as for myOLS, except
# * XF is a vector of length n that contain the covariate values (categorical or "factor")
# Outputs: same as for myOLS
}
```

Compare the results with those produced by `lm()` on the following simulated data:

```
n = 30; set.seed(0);
XF = rep(c("A","B","C"),each=10)
Y = rnorm(n) + rep(c(1,2,3),each=10)
fit1 = lm(Y ~ XF); summary(fit1) # with an intercept
fit0 = lm(Y ~ -1 + XF); summary(fit0) # without an intercept
```

**Typed Problem 4.** Exploring the equivalence of OLS and MLE in linear regression model with iid  $Normal(0, \sigma^2)$  errors.

Suppose the Nature generates the true data (pairs of  $(x_i, y_i)$ ) as shown below and then passes the vectors of  $Y$  and  $X$  values to the statistician who will then fit the simple linear regression model by ordinary least squares (OLS).

```
n = 30; set.seed(0);
X = runif(n)*4-2;
Y = 1 + 3*X + rnorm(n);
fit1 = lm(Y ~ X); summary(fit1) # beta0_hat = 1.0161; beta1_hat = 2.9304, obtained by OLS
```

4.1. Implement by hand the negative log-likelihood for the observed data.

The statistical model is  $Y_i = b_0 + b_1 x_i + \epsilon_i$ , where  $\epsilon_i$  are iid  $Normal(0, \sigma^2)$  errors;  $\sigma^2$  is the variance. The interface of the function is as follows:

```
myFullObj <- function(b,sig) {
# Inputs:
# * b is the vector of regression coefficients, b=[b0,b1];
# * sig is the standard dev of errors; sig > 0
# Output: the negative log-likelihood of the observed data (Y given X) evaluated at b and sig
}
```

Your implementation may use the built-in R function `dnorm`, but this is not required. Make sure that you implement the log-likelihood directly, rather than implementing the likelihood and then applying the log (without explicit mathematical simplifications on your end).

4.2. Assume a fixed known value for `sig` and minimize the objective function with respect to  $b$  only. E.g.,

```
sigKnown = 2; myObj1 <- function(b) {myFullObj(b,sigKnown)};
```

For optimization, specify unconstrained gradient-based search, e.g., `method="BFGS"`.

Carry out the numerical optimization of `myObj1` for several different values of `sigKnown`, examine the solutions and discuss.

Compare your minimizer(s) of `myObj1` with the `beta_OLS` solution produced by `lm()` and discuss.

4.3. Now, do not assume a known value for `sig`. Minimize `myFullObj` jointly with respect to `b` and `sig`.

Make sure that you enforce the lower bound constraint on `sig` (e.g., `sig > 10-5`) and use the gradient-based optimization routine that can handle such constraints, e.g., `"L-BFGS-B"`. To this end, follow the example from class when we studied the MLE (for iid data).

Compare your solution with that produced by `lm()` and discuss.

(Optional:) Would you expect the MLE for `sig` to be equal to the estimate produced by `lm` ("Residual standard error")? Briefly discuss.